
Circuit Safari

A Proposal for UROP/FYP/MComp Projects

Yisong Miao

yisong@comp.nus.edu.sg

Web IR / NLP Group (WING), National University of Singapore

September 23, 2025

1 Motivation




Figure 1: Comparison of word embeddings (left) and transformer circuits (right).

Word embeddings represent *meanings*. We are familiar with their geometric properties, for example $\vec{man} - \vec{woman} = \vec{king} - \vec{queen}$ (Figure 1).

Transformer circuits represent *functions*. They are sparse computational subgraphs within a transformer that approximate the full model. However, little is known about their relationships. For instance, if we combine $\text{circuit}(\text{antonymy}) + \text{circuit}(\text{coreference})$, does the combined circuit capture coreference with an emphasis on contrasting meaning? In this project, we propose constructing a “safari” of transformer circuits and developing a more precise description of circuit arithmetic.

This proposal is also available at <https://yisong.me/topics/CircuitSafari.pdf> and will remain online until the position is filled (after which we will switch to stealth mode).

2 Approach

Our work addresses two major challenges in the transformer circuits community:

Challenge 1: Limited language (and circuit) resources. Circuit discoveries have so far been made in isolation by different teams. An incomplete list includes indirect object identification (IOI), numerical comparison, subject-verb agreement (SVA), MCQ, knowledge acquisition, colored objects, extractive QA, context-free grammars, and discourse relations [WVC+23, HLV23, FCj24, LRK+23, YZX+24, OYZ+25, HPB24, MEP24, BMR+25, MWP25, MK25]. We aim to unify these efforts and extend coverage to the “classic NLP pipeline” [TDP19], filling gaps such as syntactic features, semantic features, and semantic role labeling. To this end, we will collect the necessary language resources to support circuit discovery and release the resulting circuits.

Challenge 2: Limited understanding of circuit arithmetic. Prior work has explored circuit composition [MWP25], but with limitations: (1) it focuses on synthetic probabilistic context-free grammar, detached from real-world language; (2) it examines simple set union as composition. To address (1), we will conduct systematic experiments across our Safari. To address (2), we will investigate richer arithmetic operations, such as edge subtraction or weighted set unions, and explore making composition differentiable for end-to-end optimization.

3 Organization

Who will I be working with? In this project, you will be working with WING PhD student Yisong Miao and WING lead Prof. Min-Yen Kan.

When is the starting date? January 2026 or earlier (e.g. October 2025).

What are the prerequisites? Strong natural language processing (NLP) experience, e.g., you have led a team to excel in the CS4248 / CS6101 project components (or equivalent).

What are the expected commitments and outcome? For UROP/FYP/MComp students, we expect sufficient time commitment to achieve excellence in this project. The expected outcome is a research publication, so applicants with academic inclinations are preferred.

How to get in touch? Please email Yisong and Prof. Min. We also copy Min’s message: “You **must** meet with me on Zoom or via email before deciding to take this project or any other project under my group. Please email me first to arrange a suitable timing. If you bid without consulting me, I will refuse your project bid. Read my notes on projects (<http://www.comp.nus.edu.sg/~kanmy/hyp.html>) and learn about the host organization, WING <http://wing.comp.nus.edu.sg/>, and read the requirements about entry to WING: <http://bit.ly/new-to-wing> (there’s a narrated video at <http://bit.ly/new-to-wing-yt> as well).” This project is primarily supervised by Yisong, a doctoral student under Min. Min will be overseeing the project, but not necessarily working directly with you.

References

- [BMR⁺25] Samyadeep Basu, Vlad I Morariu, Ryan A. Rossi, Nanxuan Zhao, Zichao Wang, Soheil Feizi, and Varun Manjunatha. On mechanistic circuits for extractive question-answering. In *Second Conference on Language Modeling*, 2025.
- [FCj24] Javier Ferrando and Marta R. Costa-jussà. On the similarity of circuits across languages: a case study on the subject-verb agreement task. In Yaser Al-Onaizan, Mohit Bansal, and Yun-Nung Chen, editors, *Findings of the Association for Computational Linguistics: EMNLP 2024*, pages 10115–10125, Miami, Florida, USA, November 2024. Association for Computational Linguistics.
- [HLV23] Michael Hanna, Ollie Liu, and Alexandre Variengien. How does GPT-2 compute greater-than?: Interpreting mathematical abilities in a pre-trained language model. In *Thirty-seventh Conference on Neural Information Processing Systems*, 2023.
- [HPB24] Michael Hanna, Sandro Pezzelle, and Yonatan Belinkov. Have faith in faithfulness: Going beyond circuit overlap when finding model mechanisms. In *First Conference on Language Modeling*, 2024.
- [LRK⁺23] Tom Lieberum, Matthew Rahtz, János Kramár, Neel Nanda, Geoffrey Irving, Rohin Shah, and Vladimir Mikulik. Does circuit analysis interpretability scale? evidence from multiple choice capabilities in chinchilla, 2023.
- [MEP24] Jack Merullo, Carsten Eickhoff, and Ellie Pavlick. Circuit component reuse across tasks in transformer language models. In *The Twelfth International Conference on Learning Representations*, 2024.
- [MK25] Yisong Miao and Min-Yen Kan. Discursive circuits: How do language models understand discourse relations? In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, Suzhou, China, 2025. Association for Computational Linguistics.

- [MWP25] Philipp Mondorf, Sondre Wold, and Barbara Plank. Circuit compositions: Exploring modular structures in transformer-based language models. In Wanxiang Che, Joyce Nabende, Ekaterina Shutova, and Mohammad Taher Pilehvar, editors, *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 14934–14955, Vienna, Austria, July 2025. Association for Computational Linguistics.
- [OYZ⁺25] Yixin Ou, Yunzhi Yao, Ningyu Zhang, Hui Jin, Jiacheng Sun, Shumin Deng, Zhenguo Li, and Huajun Chen. How do LLMs acquire new knowledge? a knowledge circuits perspective on continual pre-training. In Wanxiang Che, Joyce Nabende, Ekaterina Shutova, and Mohammad Taher Pilehvar, editors, *Findings of the Association for Computational Linguistics: ACL 2025*, pages 19889–19913, Vienna, Austria, July 2025. Association for Computational Linguistics.
- [TDP19] Ian Tenney, Dipanjan Das, and Ellie Pavlick. BERT rediscovers the classical NLP pipeline. In Anna Korhonen, David Traum, and Lluís Màrquez, editors, *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 4593–4601, Florence, Italy, July 2019. Association for Computational Linguistics.
- [WVC⁺23] Kevin Ro Wang, Alexandre Variengien, Arthur Conmy, Buck Shlegeris, and Jacob Steinhardt. Interpretability in the wild: a circuit for indirect object identification in GPT-2 small. In *The Eleventh International Conference on Learning Representations*, 2023.
- [YZX⁺24] Yunzhi Yao, Ningyu Zhang, Zekun Xi, Mengru Wang, Ziwen Xu, Shumin Deng, and Huajun Chen. Knowledge circuits in pretrained transformers. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024.